

# Article Duplicates

Brian Lynnerup Pedersen



Kongens Lyngby 2014

Technical University of Denmark  
Department of Applied Mathematics and Computer Science  
Matematiktorvet, building 303B,  
DK-2800 Kgs. Lyngby, Denmark  
Phone +45 4525 3351  
[compute@compute.dtu.dk](mailto:compute@compute.dtu.dk)  
[www.compute.dtu.dk](http://www.compute.dtu.dk)

# Summary (English)

---

The goal of this thesis is create is to document my work with implementing a working prototype of using algorithms to find article duplicates in a large corpus of articles.



# Summary (Danish)

---

Målet for denne afhandling er at dokumentere mit arbejde med at finde artikel duplikater, i et større artikel corpus, ved brug af algoritmer.



# Preface

---

This thesis was prepared at the department of Informatics and Mathematical Modelling at the Technical University of Denmark in fulfilment of the requirements for acquiring an B.Eng. in Informatics.

The project is equal to 20 ECTS points.

**This thesis is protected by confidentiality. No information from this thesis can be handed off to any party without signed permission.**

The thesis deals with the issue of finding articles that are duplicates in a large corpus of articles. This is done using various algorithms and is implemented in C#.

The thesis consists of ...

Lyngby, 09-June-2014

*Not Real*





# Acknowledgements

---

I would like to thank my supervisors from DTU, Inge Li Gørtz and Philip Bille, for the help they provided to my project. Also I would like to thank my company Infomedia, for letting me do my project with them, in particular my project leader Klaus Wenzel Jørgensen and Rene Madsen.



# Contents

---

<b>Summary (English)</b>	<b>i</b>
<b>Summary (Danish)</b>	<b>iii</b>
<b>Preface</b>	<b>v</b>
<b>Acknowledgements</b>	<b>vii</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 General - Terms and Rules</b>	<b>3</b>
<b>3 Algorithms in General</b>	<b>5</b>
3.1 Initial Considerations . . . . .	5
3.2 Algorithms Used . . . . .	6
3.3 Optimizing Performance . . . . .	7
3.3.1 Stop Words . . . . .	7
3.3.2 Stemming . . . . .	8
3.4 Semaphore Tagging . . . . .	9
3.4.1 Cons to Semaphore . . . . .	9
<b>4 Excel Test</b>	<b>11</b>
<b>A Appendix</b>	<b>13</b>



## CHAPTER 1

# Introduction

---

I have been working for Infomedia<sup>1</sup> since my fourth semester at DTU (march 2012), while studying as an IT engineer (B.Eng.). Infomedia is in short a company that deals with news monitoring.

Infomedia is the result of a fusion between Berlingske Avisdata and Polinfo in 2002, which means that Infomedia is partly owned by JP/Politikens Hus<sup>2</sup> and Berlingske Media<sup>3</sup>. It is a company with around 130 employees, of which a fair amount is student aides like myself. Infomedia has various departments, which includes an economy, sales, analysis and an IT department amongst others.

I am employed in the IT department as a student programmer.

Infomedia deals with news monitoring, which means that we have an inflow of articles<sup>4</sup> from various newspapers, news sites, television and radio media, which we then monitor for content that is of interest to our clients. This can be a client that wishes to know when their firm is mentioned in the press or a product they are using, if that is being mentioned. Infomedia have also begun monitoring social media. Infomedia then sells various solutions to clients, for them to get this news monitoring.

---

<sup>1</sup>[www.infomedia.dk](http://www.infomedia.dk)

<sup>2</sup>[www.jppol.dk](http://www.jppol.dk)

<sup>3</sup>[www.berlingskemedias.dk](http://www.berlingskemedias.dk)

<sup>4</sup>Articles are sent to Infomedia daily, this can be more than 40,000 articles per day.

One of the things that Infomedia tries to do, is that we want to present our clients with a fast overview of the articles in which terms<sup>5</sup>, that trigger our news monitoring, appear. Many local newspapers are today owned by bigger media houses (like the owners of Infomedia) and as such, they will feature a lot of the articles that have also been printed in the "mother paper". This will make the same (or roughly the same<sup>6</sup>) article appear many times in news monitoring. In an effort to make the list of articles presented to the clients, easy to look at, and preventing a client having to read the "same" article many times, Infomedia has a wish to cluster article duplicates. Infomedia can then present the client with a list of articles and in that list have further sub lists that contains duplicates of the original article<sup>7</sup>. This also have an economic factor as clients are charged per article read.

Another issue, is the issue of copyrights and when the same article will appear in different media, but without content given from the author of that article. An example that is often happening is that news telegrams from Reuters<sup>8</sup> or Ritzau<sup>9</sup> is published in a newspaper, but without the source indication. All news media are of course interested in knowing when their material is being published in competing media. This how ever can be tricky business, as official rules on the matter is incredible fuzzy.

I will in this thesis try and look into various ways of identifying article duplicates within a test corpus<sup>10</sup> of articles, by using algorithms. The long term goal for Infomedia is having this being implemented in the inflow of articles, and having a look back functionality so that we can group duplicates not just for one day, but for a longer period of time.

---

<sup>5</sup>A term is, in short, a word or a combination of words. For the rest of my thesis a term will however only be a single word.

<sup>6</sup>Articles can be slightly edited in order to make them fit into the layout of the various papers.

<sup>7</sup>Or the longest article rather, as this will tend to contain the most information.

<sup>8</sup>[www.reuters.com](http://www.reuters.com)

<sup>9</sup>[www.ritzau.dk](http://www.ritzau.dk)

<sup>10</sup>A days worth of articles from 10/31/2013 - totalling 22.787 articles.

## CHAPTER 2

# General - Terms and Rules

---

As there is a lot of terms used in this thesis, a short introduction to the most used are in order.

- **Article:** For this thesis, a digital document containing the contents of a piece of news. Could originate from papers, magazines, TV or other forms of media. For this thesis a document corresponds to an article. Articles are in their electronic form stored at Infomedia as XML files, I will throughout this thesis only deal with the part of the XML files that contains data of value to me in this assignment. This being *Tags* (see below), *Headline*, *Sub headline* and *Article Text*.
- **Corpus:** From Latin meaning *body*. In this thesis that describes the test set of articles being used a test set throughout my thesis.
- **Tag:** Used in Ontology<sup>1</sup> to create words that describes the contents of an article.
- **Term:** Basically a word. A term will be something that can be searched for.

---

<sup>1</sup>[http://en.wikipedia.org/wiki/Ontology\\_\(information\\_science\)](http://en.wikipedia.org/wiki/Ontology_(information_science))

Hvornår er en artikel et duplikat (korte artikler (breaking news), hvornår er to artikler "tilstrækkeligt" forskellige?) blah about copyright rules in Denmark...



## CHAPTER 3

# Algorithms in General

---

Before any sort of work can be done, one must consider various algorithm to work with. There are several text matching algorithms available for free on the Internet, and if one has the money for it, there are companies that can develop a specialized algorithm for you. As I do not have a lot of money (and paying for someone else, to do an algorithm for me, kind of defeats the purpose of this whole thesis) I have gone with the first option and found a free basic algorithm on the Internet, as well as contemplated to create my own algorithm from scratch.

## 3.1 Initial Considerations

Infomedia already have an algorithm implemented in the inflow to make a rough comparison of the articles coming in. However, the thought is that a combination of several algorithms would provide a better and more granular view of the articles as they are being compared. A new algorithm should be one that was specialized in text matching. It should also be an algorithm that would work in different manner than what Infomedia already have implemented<sup>1</sup>, as having

---

<sup>1</sup>More on the algorithm already implemented in the next section.

two algorithms that work in more or less the same fashion would not produce results of much interest.

As the current implementation is rather fast, it could prove useful to have the algorithm that is already implemented, to do the initial split and then have a slower (but more thorough) algorithm look at the *interesting* article comparisons. Initially I have looked at two algorithms to fill this need, *Longest Common Substring* and *Semaphore Tag Matching* - an algorithm I would make from scratch.

## 3.2 Algorithms Used

**Term Frequency - Inverse Document Frequency<sup>2</sup> (Cosine)**, generates a vector from each document. **I will not cover this algorithm in detail, as it is not the focus for my thesis.** This is the algorithm already implemented in the inflow today (a version of it). Each word in every article is added to a *word map* which contains all the words of all articles in the corpus being checked (which also is used to create the document vector). The word map is used to generate a weight of each word (a word occurring in many articles will have less weight than a word that is only present in a few articles). Each word generates a bit of the articles total vector, a word that occurs in all documents will have a very short vector, a more rare word will have a longer (and therefore weigh heavier) in the article vector.

Once the word map is created, the articles are then scored based on the words in the article and the word map. This score is a vector, which is used to compare the article with other articles. This is done one article at a time (although done with parallel coding to speed up the process). As the word map is generated each time the algorithm is run, the word map can (and probably will) differ from each run (if the corpus of articles are being changed). Infomedia is therefore talking about implementing this bit differently, and building a constant word map, that only gets updated with each run, not overwritten.

The algorithm then returns a list of article comparisons (based on a threshold set by the user), with article ID<sup>3</sup> and scores. So each comparison has the ID of "*article 1*" and "*article 2*" and their score (the angle between the two vectors, in a multidimensional universe). The closer the score is to the value 1.0 the more similar are two articles. A score of 0.0 indicates that two articles has nothing in common (according to this algorithm, I will discuss this point in the

---

<sup>2</sup><http://en.wikipedia.org/wiki/Tf-idf>

<sup>3</sup>Each article has their own unique ID.

next paragraph), whereas a score of 1.0 indicates two perfectly identical articles (according to this algorithm). This algorithm and it has a good O-notation (!!!  $N \log N$ ???? !!!).

**Problem:** This string comparison is very sensitive to short articles, for instance "A man walks his dog in the park" and "A dog walks his man in the park" would result in returning a cosine value of 1.0, due to the term vector. This problem is very unlikely to yield false positives.

**Longest Common Substring (LCS)**, compares documents in pairs as well. A general implementation would be to have a list of documents and then compare a document to every other document in the list. The algorithm will then return the length of the longest common substring. By default<sup>4</sup> the LCS algorithm checks the contents of a string character by character. When initialized the algorithm creates an double array, each time a match is found (when two identical characters are found ('a' and 'a' for instance)) the algorithm marks that in the array by adding a number.

!!!!!!!!!! BESKRIV LCS I DETALJER HER - BILLEDER!!!!!!

I will use the basic implementation of this algorithm in my thesis, I will then try and modify it to work better in context with finding article duplicates, instead of just the longest common substring.

**Problem:** This algorithm is very prone to fail in cases where there have been made alterations to the article in question. A word change in the middle of one of two otherwise identical articles will result in a 50 % match. If the article in question have been obfuscated with many changed words, the LCS will be extremely short. This is a high risk problem, as article duplication will often involve changing words.

## 3.3 Optimizing Performance

### 3.3.1 Stop Words

Stop word<sup>5</sup> removal would improve running time (performance) of both algorithms, and would pose little threat of causing either algorithm to fail (finding

<sup>4</sup>[http://en.wikibooks.org/wiki/Algorithm\\_Implementation/Strings/Longest\\_common\\_substring](http://en.wikibooks.org/wiki/Algorithm_Implementation/Strings/Longest_common_substring)

<sup>5</sup>[http://en.wikipedia.org/wiki/Stop\\_words](http://en.wikipedia.org/wiki/Stop_words)

false positives). The exception to this could be very short articles (like breaking news articles<sup>6</sup>), that only contains common words, like "Man walks away". Depending on the stop word list, this article could end up being *null*<sup>7</sup>. We can safely (with respects to the previously addressed problem) remove stop words, as they don't provide any *semantic*<sup>8</sup> value to the text.

#### 3.3.1.1 Cosine

For the cosine algorithm removal of stop words would improve performance, by reducing the size of the *Magnitude Vector*. However, as the algorithm already weighs words based on their number of occurrence, and stop words by nature occur often, this will have a minimal impact on performance.

#### 3.3.1.2 LCS

In regards to the LCS algorithm, the performance would also be improved, as the algorithm wouldn't need to traverse as many characters. This would reduce the length of the longest common substring, but it would be unlikely that it would affect the outcome of the algorithm, as stop words are rarely changed when duplicating articles.

### 3.3.2 Stemming

*Stemming*<sup>9</sup> would improve performance of the cosine algorithm. This would not affect accuracy as stemming does not ruin the semantic content.

#### 3.3.2.1 Cosine

Stemming improves the performance of the cosine algorithm as this will reduce the number of words in the *vector space*. As words would be reduced to their

---

<sup>6</sup>Breaking News articles is a thing of the 2000s. With the spread of media onto the Internet, an article is no longer a static printed piece of news in a paper. News can be published instantly on the web, and then updated as information are received. However, Breaking News articles can still be printed in the paper article.

<sup>7</sup>An article with no text data.

<sup>8</sup><http://en.wikipedia.org/wiki/Semantics>

<sup>9</sup><http://en.wikipedia.org/wiki/Stemming>

base form (Danish: *Grundform*), words used several times, but with different endings would be counted as the same word. When using weighed evaluation this would actually improve performance to some extend.

### 3.3.2.2 LCS

Stemming would not improve the performance of the LCS algorithm by much. As this algorithm matches characters one by one, it would make little difference if the words are stemmed or not. For LCS stemming could actually end up decreasing performance, as this operation would consume time while preparing the files, with no (or little) impact to the performance of the algorithm.

## 3.4 Semaphore Tagging

Another way of finding duplicates is by looking at each articles semaphore tags. When Infomedia receives articles in the inflow, these articles are enriched with *Semaphore Tags*<sup>10</sup>. Each article then gets a number of tags based on what terms are found in the article. A way of finding article duplicates could be through creating an algorithm that would check a pair of articles with their respective semaphore tags.

### 3.4.1 Cons to Semaphore

As these tags are more a general indicator to an articles contents than an actual text matching algorithm, it will provide very little value as a stand alone implementation. Each article will only contain tags for the terms worth of note. If 100 articles all contained various doings of a minister (picking up the children, going to meetings, being involved in a crisis) many of the same tags would be present in the 100 articles, this could be the ministers name, political party and other general tags that are linked to this minister. It would therefore be hard to decipher much information that could truly and uniquely link two articles together just by doing this. However this could be used to enhance another algorithm (for instance one of the two mentioned above). I will therefore not look any further into implementing this algorithm in this thesis.

---

<sup>10</sup>Semaphore tags are tags that describes the contents of the article, an article about financial fraud would contain the tag *Economic Crime*. These tags are created by the Infomedia Ontology<sup>11</sup> team. They are creating rules for when a certain word (or words) appear in certain context, then an article will be tagged with a certain semaphore tag.

Belowthis should be a picture!




Figure 3.1: TEXT!

## CHAPTER 4

# Excel Test

---

Insert Excel data somehow...





## APPENDIX A

# Appendix

---

This appendix is full of stuff ... add more

Document from Infomedia about the general Vector Spaced Search.

Thesis document from Inge

Links to Wikipedia

The algorithm course book..

Example XML document (article) - to show what the XML documents look like.

